# Jongse Park                                                        April 2024

| | |
|---|---|
| **Contact Information** | School of Computing<br>KAIST<br>291 Daehak-ro, Yuseong-gu      *E-mail:* jspark@casys.kaist.ac.kr<br>Daejeon, South Korea, 34141      *URL:* https://jongse-park.github.io |

**Research Interests**

Computer Architecture, HW/SW Codesign, AI Systems, Autonomous Systems

**Employment**

| | |
|---|---|
| **Associate Professor**. **KAIST** | Mar. 2024–date |
| **Assistant Professor**. **KAIST** | Dec. 2019–Feb. 2024 |
| **Product Engineer**. **Bigstream Solutions Inc.** | Jun. 2018–Nov. 2019 |

**Education**

**Ph.D. in Computer Science. Georgia Institute of Technology**      Aug. 2013–Aug. 2018
- Advisor: Prof. Hadi Esmaeilzadeh
- Dissertation: *Breaking the Abstractions for Productivity and Performance in the Era of Specialization*

**M.S. in Computer Science. KAIST**      Feb. 2012
- Advisor: Prof. Seungryoul Maeng
- Thesis: *Dynamic Resource Reconfiguration on the Cloud for Improving Data Locality*

**B.E. in Computer Science and Engineering. Sogang University**      Feb. 2010
- Graduated with Honors

**Honors and Awards**

| | |
|---|---|
| Distinguished paper award. IEEE Symposium on High Performance Computer Architecture. "TABLA: A Unified Template-Based Framework for Accelerating Statistical Machine Learning" | 2016 |
| Honorable Mention in IEEE Micro Top Picks from 2014 Computer Architecture Conferences. "General-Purpose Code Acceleration with Limited-Precision Analog Computation" | 2015 |
| Kwanjeong Foundation Scholarship, Kwanjeong Educational Foundation (KEF) | 2013–2018 |
| National Full Scholarship, KAIST | 2010–2012 |
| Dean's Honored Graduate, Ranked $3^{rd}$ among graduates of the class of 2010 | 2010 |
| DMC General Management Track Scholarship, Samsung Electronics Co., Ltd | 2009 |
| Academic Scholarship, Sogang University, 7 semesters | 2004–2009 |

**Refereed Conference Papers**

1. M. Kim, J. Hwang, G. Heo, S. Cho, D. Mahajan, **J. Park**, "Accelerating String-key Learned Index Structures via Memoization-based Incremental Training" in *International Conference on Very Large Data Bases (**VLDB**)*, August 2024 (To Appear).

2. Y. Kim, C. Oh, J. Hwang, W. Kim, S. Oh, Y. Lee, H. Sharma, A. Yazdanbakhsh, **J. Park**, "DaCapo: Accelerating Continuous Learning in Autonomous Systems for Video Analytics" in *International Symposium on Computer Architecture (**ISCA**)*, June 2024 (To Appear).

3. G. Heo, S. Lee, J. Cho, H. Choi, S. Lee, H. Ham, G. Kim, D. Mahajan, **J. Park**, "NeuPIMs: NPU-PIM Heterogeneous Acceleration for Batched LLM Inferencing" in *International Conference on Architectural Support for Programming Languages and Operating Systems (**ASPLOS**)*, April 2024 (To Appear).

4. S. Ghodrati, S. Kinzer, H. Xu, R. Mahapatra, Y. Kim, B. H. Ahn, D. K. Wang, L. Karthikeyan, A. Yazdanbakhsh, **J. Park**, N. S. Kim, H. Esmaeilzadeh, "Tandem Processor: Grappling with Emerging Operators in Neural Networks" in *International Conference on Architectural Support for Programming Languages and Operating Systems (**ASPLOS**)*, April 2024 (To Appear).

5. Sunho Lee, Seonjin Na, Jungwoo Kim, Jongse Park, and Jaehyuk Huh, "Tunable Memory Protection for Secure Neural Processing Units" in *The 40th IEEE International Conference on Computer Design (**ICCD**)*, October 2022.

6. Bokyeong Kim, Soojin Hwang, Sanghoon Cha, Chang Hyun Park, Jongse Park, and Jaehyuk Huh, "Supporting Dynamic Translation Granularity for Hybrid Memory Systems" in *The 40th IEEE International Conference on Computer Design (**ICCD**)*, October 2022.

7. Joon Kyung Kim, Byung Hoon Ahn, Sean Kinzer, Soroush Ghodrati, Rohan Mahapatra, Brahmendra Yatham, Dohee Kim, Parisa Sarikhani, Babak Mahmoudi, Divya Mahajan, Jongse Park, Hadi Esmaeilzadeh, "Yin-Yang: Programming Abstraction for Cross-Domain Multi-Acceleration", in ***IEEE Micro**, special issue on Compiling for Accelerators*, 2022.

8. Jinwoo Hwang, Minsu Kim, Daeun Kim, Seungho Nam, Yoonsung Kim, Dohee Kim, Hardik Sharma, Jongse Park, "CoVA: Exploiting Compressed-Domain Analysis to Accelerate Video Analytics", in *USENIX Annual Technical Conference (**ATC**)*, July 2022.

9. Seungbeom Choi, Sunho Lee, Yeonjae Kim, Jongse Park, Youngjin Kwon, and Jaehyuk Huh, "Serving Heterogeneous Machine Learning Models on Multi-GPU Servers with Spatio-Temporal Sharing", in *USENIX Annual Technical Conference (**ATC**)*, July 2022.

10. S. Lee, J. Kim, S. Na, **J. Park**, and J. Huh, "TNPU: Supporting Trusted Execution with Tree-less Integrity Protection for Neural Processing Unit" in *The 27th IEEE International Symposium on High-Performance Computer Architecture (**HPCA**)*, February 2022. [To appear]

11. S. Na, S. Lee, Y. Kim, **J. Park**, and J. Huh, "Common Counters: Compressed Encryption Counters for Secure GPU Memory" in *The 27th IEEE International Symposium on High-Performance Computer Architecture (**HPCA**)*, February 2021.

12. S. Ghodrati, H. Sharma, S. Kinzer, A. Yazdanbakhsh, **J. Park**, N. Kim, D. Burger, and H. Esmaeilzadeh, "Mixed-Signal Charge-Domain Acceleration of Deep Neural Networks through Interleaved Bit-Partitioned Arithmetic" in *The 29th International Conference on Parallel Architectures and Compilation Techniques (**PACT**)*, October 2020.

13. Y. Li, **J. Park**, M. Alian, Y. Yuan, Q. Zheng, P. Pan, R. Wang, A. Schwing, H. Esmaeilzadeh, N. Kim, "A Network-Centric Hardware/Algorithm Co-Design to Accelerate Distributed Training of Deep Neural Networks," *The 50th Annual IEEE/ACM International Symposium on Microarchitecture (**MICRO**)*, October 2018.

14. H. Sharma, **J. Park**, B. Samynathan, B. Robatmili, S. Mirkhani, H. Esmaeilzadeh, "From Tensors to FPGAs: Accelerating Deep Learning," *A Symposium on High Performance Chips (**Hot Chips**)*, August 2018.

15. H. Sharma, **J. Park**, N. Suda, L. Lai, B. Chau, J. Kim, V. Chandra, H. Esmaeilzadeh, "Bit Fusion: Bit-Level Dynamically Composable Architecture for Accelerating Deep Neural Networks," *International Symposium on Computer Architecture (**ISCA**)*, June 2018.

16. **J. Park**, H. Sharma, D. Mahajan, J. Kim, P. Olds, H. Esmaeilzadeh, "Scale-Out Acceleration for Machine Learning," in *The 50th Annual IEEE/ACM International Symposium on Microarchitecture (**MICRO**)*, October 2017.

17. **J. Park**, E. Amaro, D. Mahajan, B. Thwaites, H. Esmaeilzadeh, "AxGames: Towards Crowdsourcing Quality Target Determination in Approximate Computing," in *International Conference on Architectural Support for Programming Languages and Operating Systems (**ASPLOS**)*, April 2016.

18. H. Sharma, **J. Park**, D. Mahajan, E. Amaro, J. Kim, C. Shao, A. Mishra, H. Esmaeilzadeh "From High-Level Deep Neural Models to FPGAs," in *The 49th Annual IEEE/ACM International Symposium on Microarchitecture (**MICRO**)*, October 2016.

19. D. Mahajan, **J. Park**, E. Amaro, H. Sharma, A. Yazdanbaksh, J. Kim, H. Esmaeilzadeh, "Tabla: A Unified Template-based Framework for Accelerating Statistical Machine Learning," in *The 22nd IEEE Symposium on High Performance Computer Architecture (**HPCA**)*, March 2016.

    ***(Distinguished Paper Award)***

20. D. Mahajan, A. Yazdanbaksh, **J. Park**, B. Thwaites, H. Esmaeilzadeh, "Towards Statistical Guarantees in Controlling Quality Tradeoffs in Approximate Acceleration," in *International Symposium on Computer Architecture (**ISCA**)*, June 2016.

21. A. Yazdanbakhsh, **J. Park**, H. Sharma, P. Lotfi-Kamran, H. Esmaeilzadeh, "Neural Acceleration for GPU Throughput Processors," in *The 48th Annual IEEE/ACM International Symposium on Microarchitecture (**MICRO**)*, December 2015.

22. **J. Park**, H. Esmaeilzadeh, X. Zhang, M. Naik, W. Harris, "FLEXJAVA: Language Support for Safe and Modular Approximate Programming," in *The 10th Joint Meeting of the European Software Engineering Conference and the ACM SIGSOFT Symposium on the Foundations of Software Engineering (**FSE**)*, September 2015.

23. A. Yazdanbakhsh, D. Mahajan, B. Thwaites, **J. Park**, A. Nagendrakumar, S. Sethuraman, K. Ramkrishnan, N. Ravindran, R. Jariwala, A. Rahimi, H. Esmailzadeh, K. Bazargan, "AXILOG: Language Support for Approximate Hardware Design," in *Design Automation and Test in Europe (**DATE**)*, March 2015.

24. R. S. Amant, A. Yazdanbakhsh, **J. Park**, B. Thwaites, H. Esmaeilzadeh, A. Hassibi, L. Ceze, D. Burger, "General-Purpose Code Acceleration with Limited-Precision Analog Computation," in *The 41th International Symposium on Computer Architecture (**ISCA**)*, June 2014.

    ***(Nominated for CACM Research Highlights; Honorable Mention in IEEE Micro Top Picks)***

25. B. Thwaites, G. Pekhimenko, A. Yazdanbakhsh, **J. Park**, G. Mururu, H. Esmaeilzadeh, O. Mutlu, T. Mowry, "Rollback-Free Value Prediction with Approximate Loads," in *The 24th International Conference on Parallel Architectures and Compilation Techniques (**PACT**)*, August 2014.

26. J. Choi, **J. Park**, J. Seol, and S. Maeng, "Isolated Mini-domain for Trusted Cloud Computing," in *The 13th International Symposium on Cluster, Cloud, and Grid Computing (**CCGrid**)*, May 2013.

27. **J. Park**, D. Lee, B. Kim, J. Huh, S. Maeng, "Locality-aware Dynamic VM Reconfiguration on MapReduce Clouds," in *The 21st International ACM Symposium on High-Performance Parallel and Distributed Computing (**HPDC**)*, June 2012.

**Refereed Journal Articles**

1. J. Park, S. Kang, S. Lee, T. Kim, **J. Park**, Y. Kwon, and J. Huh, "Hardware Hardened Sandbox Enclaves for Trusted Serverless Computing" in *IEEE Transactions on Architecture and Code Optimization (**TACO**)*, 2023.

2. S. Noh, J. Koo, S. Lee, **J. Park**, and J. Kung, "FlexBlock: A Flexible DNN Training Accelerator with Multi-Mode Block Floating Point Support" in *IEEE Transactions on Computers (**TC**)*, 2023.

3. S. Lee, R. Hwang, **J. Park**, and M. Rhu, "HAMMER: Hardware-friendly Approximate Computing for Self-attention with Mean-redistribution and Linearization" in *IEEE Computer Architecture Letters (**CAL**)*, 2023.

4. W. Seo, S. Cha, Y. Kim, J. Huh, and **J. Park**, " SLO-aware Inference Scheduler for Heterogeneous Processors in Edge Platforms" in *Transactions on Architecture and Code Optimization (**TACO**)*, 2021.

5. D. Mahajan, K. Ramkrishnan, R. Jariwala, A. Yazdanbakhsh, **J. Park**, B. Thwaites, A. Nagendrakumar, A. Rahimi, H. Esmaeilzadeh, K. Bazargan, "AXILOG: Abstractions for Approximate Hardware Design and Reuse," in ***IEEE Micro**, special issue on Alternative Computing Designs and Technologies*, October 2015.

**Refereed Workshop Papers**

1. Y. Lee, **J. Park**, "LVS: A Learned Video Storage for Fast and Efficient Video Understanding" in *Efficient Deep Learning for Computer Vision (**ECV**) in conjunction with CVPR*, June 2024 (To Appear).

2. H. Sharma, **J. Park**, E. Amaro, B. Thwaites, P. Kotha, A. Gupta, J. Kim, A. Mishra, H. Esmaeilzadeh, "DNNWEAVER: From High-Level Deep Network Models to FPGA Acceleration," in *The Second Workshop on Cognitive Architectures (**CogArch**) in conjunction with ASPLOS*, April 2016.

3. D. Mahajan, A. Yazdanbakhsh, **J. Park**, B. Thwaites, H. Esmaeilzadeh, "Prediction-Based Quality Control for Approximate Accelerators," in *The Second Workshop on Approximate Computing Across the System Stack (**WACAS**) in conjunction with ASPLOS*, March 2015.

4. **J. Park**, K. Ni, X. Zhang, H. Esmaeilzadeh, M. Naik, "Expectation-Oriented Framework for Automating Approximate Programming,", in *The First Workshop on Approximate Computing Across the System Stack (**WACAS**) in conjunction with ASPLOS*, March 2014.

5. A. Yazdanbakhsh, B. Thwaites, **J. Park**, H. Esmaeilzadeh, "Methodical Approximate Hardware Design and Reuse," in *The First Workshop on Approximate Computing Across the System Stack (**WACAS**) in conjunction with ASPLOS*, March 2014.

6. A. Yazdanbakhsh, R. Amant, B. Thwaites, **J. Park**, H. Esmaeilzadeh, A. Hassibi, L. Ceze, D. Burger, "Toward General-Purpose Code Acceleration with Analog Computation," in *The First Workshop on Approximate Computing Across the System Stack (**WACAS**) in conjunction with ASPLOS*, March 2014.

7. B. Thwaites, A. Yazdanbakhsh, **J. Park**, H. Esmaeilzadeh, "Bio-Accelerators: Bridging Biology and Silicon for General-Purpose Computing," in *Wild and Crazy Ideas (**WACI**) in conjunction with ASPLOS*, March 2014.

| | | |
|---|---|---|
| **Research Experience** | **Research Assistant**. Alternative Computing Technology (ACT) Lab | Aug. 2013–Aug. 2018 |
| | • Georgia Institute of Technology | |
| | • Advisor: Prof. Hadi Esmaeilzadeh | |
| | **Visiting Researcher**. Alternative Computing Technology (ACT) Lab | Jan. 2018–Aug. 2018 |
| | • University of California, San Diego | |
| | • Advisor: Prof. Hadi Esmaeilzadeh | |
| | **Research Intern**. Architecture Research Group (ARG) | May 2017–Aug. 2017 |
| | • NVIDIA Research | |
| | • Mentors: Dr. Arslan Zulfiqar and Dr. Eiman Ebrahimi | |
| | • Manager: Dr. Stephen Keckler | |
| | **Research Intern**. Catapult team | Jan. 2016–May 2016 |
| | • Microsoft Research | |
| | • Mentor: Dr. Eric Chung | |
| | • Manager: Dr. Doug Burger | |
| | **Research Assistant**. Computer Architecture (CA) Lab | Feb. 2010–Jul. 2013 |
| | • Korea Advanced Institute of Science and Technology (KAIST) | |
| | • Advisor: Prof. Seungryoul Maeng | |

| | | |
|---|---|---|
| **Teaching Experience** | **Instructor.** | |
| | • CS610: Parallel Processing | Spring 2024 |
| | • CS311: Computer Organization | Spring 2024 |
| | • CS411: System for AI | Fall 2023 |
| | • CS510: Computer Architecture | Spring 2023 |
| | • CS230: System Programming | Fall 2022 |
| | • CS311: Computer Organization | Spring 2022 |
| | • CS230: System Programming | Fall 2021 |
| | • CS492: Special Topic in Computer Science: System for Artificial Intelligence | Spring 2021 |
| | • CS230: System Programming | Fall 2020 |
| | • CS492: Special Topic in Computer Science: System for Machine Learning | Spring 2020 |
| | **Teaching Assistant.** | |

|  |  |  |  |  |
|---|---|---|---|---|
| • CS3220: | Processor Design | Georgia Institute of Technology | Fall 2016 |
| • CS3220: | Processor Design | Georgia Institute of Technology | Fall 2014 |
| • CS8803: | Alternative Computing Technology | Georgia Institute of Technology | Spring 2014 |
| • CS211: | Digital System and Lab. | KAIST | Spring 2011 |
| • CS311: | Embedded Computer Systems. | KAIST | Fall 2010 |

**Technical Skills**  Programming languages: C/C++, Java, Python, CUDA, Verilog, Bash, JavaScript, HTML
Development Tools: Tensorflow, Amazon EC2, Spark, Hadoop, Chord, LLVM

**References Available to Contact**

**Hadi Esmaeilzadeh.** Professor, UCSD                                   hadi@eng.ucsd.edu

**Nam Sung Kim** Professor, UIUC                                          nskim@illinois.edu

**Doug Burger.** Technical Fellow and Corporate VP, Microsoft Research       dburger@microsoft.com

**Stephen W. Keckler.** VP of Architecture Research, NVIDIA Research         skeckler@nvidia.com

**Eric Chung.** VP of AI Computing, NVIDIA                                eschung@gmail.com