

PyTorchSim: A Comprehensive, Fast, and Accurate NPU Simulation Framework

Wonhyuk Yang^{*}, Yunseon Shin^{*}, Okkyun Woo^{*}, Geonwoo Park[§], Hyunkyu Ham,
Jeehoon Kang^{††}, Jongse Park^{††}, Gwangsun Kim

POSTECH

Parallel System Architecture Lab.

|| [†] FuriosaAI

^{††} KAIST





^{*} Equal contribution

[§] Currently with Samsung Electronics

AI Model & Hardware Trends

- Deep Neural Network (DNN) are growing exponentially in size
- **Hardware performance** improves relatively slowly

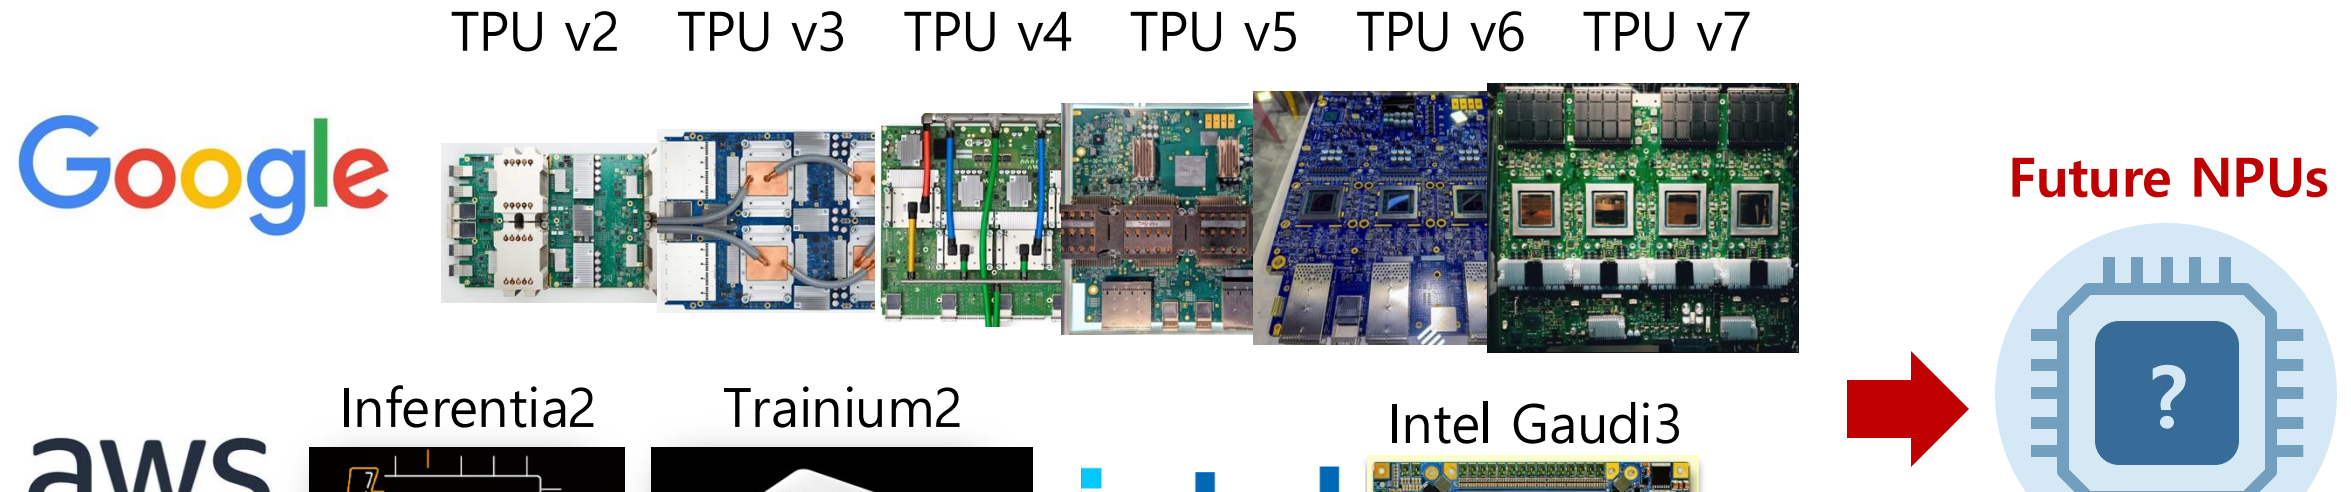


We need **high-performance** & **efficient** hardware

Evolution of Neural Processing Units

Neural Processing Units (NPU) are designed to address this challenge

- Efficient dataflow units (e.g., systolic arrays, adder-tree)




NPU simulators play a crucial role in designing **NPU**s

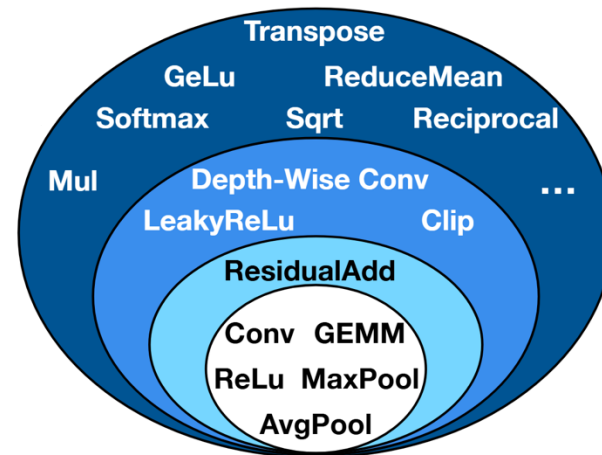
NPU Simulator Requirements – AI Model

- Both **inference** and **training** are important
- **(Vector) operations** are becoming increasingly diverse
- **Sparsity** is now widely exploited for efficiency
 - Latency varies with input data (i.e., **data-dependent timing behavior**)


Inferentia2
(Inference)



Trainium2
(Training)




Source: Tandem Processor
[Ghodrati, Soroush, et al., ASPLOS 2024]




NPU Simulator Requirements – NPU Hardware

- Need to accurately model key components of modern NPUs
- **NPU core:** matrix, vector, and scalar units and scratchpad memory
- **Shared resources:** interconnect and DRAM
 - Detailed simulation often necessary to model inter-core contention

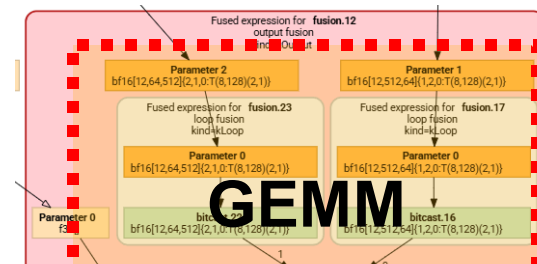


NPU Simulator Requirements – Compiler

- Necessary to bridge AI models and NPU hardware
- Lowers DNN models into machine code
- Apply various optimizations
- Enables full-model simulation for both inference and training



Operation fusion



For compiler support, an **NPU ISA** must first be defined

NPU ISA


- No de-facto standard ISA for NPUs (unlike CPUs and GPUs)
- RISC-V can be a strong candidate for an NPU ISA
 - Generality:
 - Vector-length agnostic design
 - A generic interface (VCIX) for dataflow units
 - Extensibility: Reserved opcode space for custom instructions
 - Openness: Contributions from both academia and industry
 - SW Infrastructure: Rich software ecosystem (e.g., compilers, simulators)



Existing NPU Models

- Various NPU simulators and analytical models have been proposed

Timeline



None of them meets all the key requirements

1. Full AI model support



2. AI compiler support



3. NPU core & shared resources



We propose **PyTorchSim**, which can better satisfy the key requirements for NPU simulation

Contents

- Background / Motivation
- **PyTorchSim**
 - PyTorchSim overview
 - NPU core modeling
 - Compilation flow
- Methodology & Evaluation
- Case study
 - Impact of DNN Training Hyperparameter
 - Scheduling for Chiplet-based NPUs
- Summary

PyTorchSim Framework: Overview

**PyTorch AI
model**

- Runs PyTorch models directly



 **PyTorch 2.x
NPU Backend**

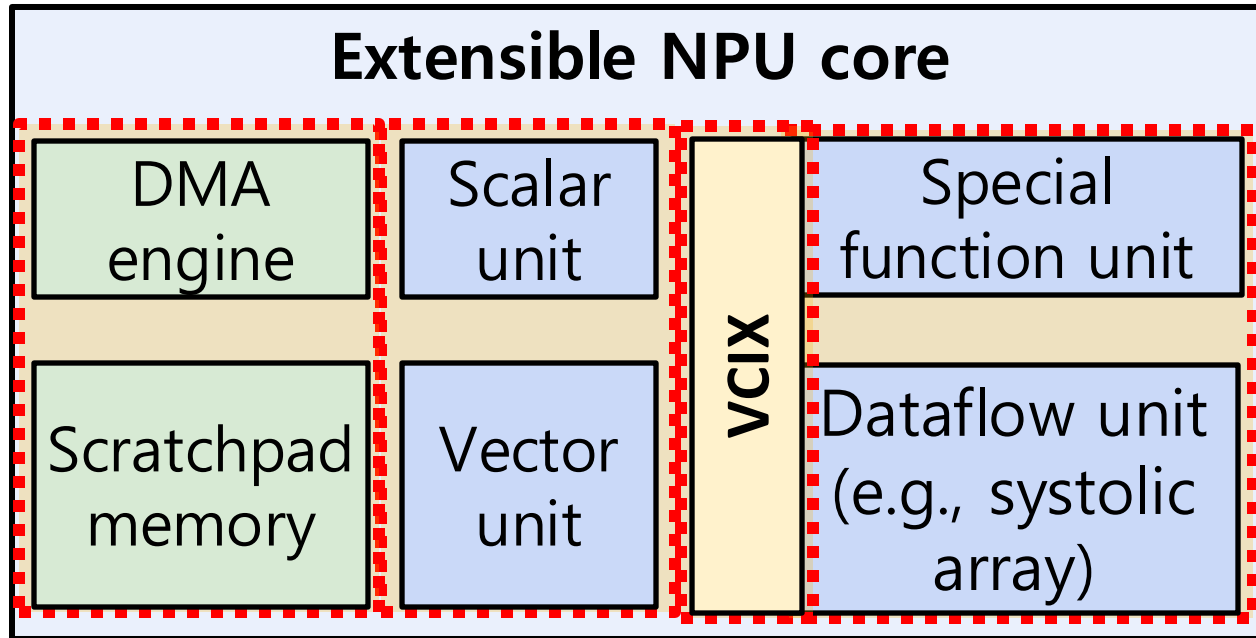
- Compiles the models to a RISC-V based NPU ISA



NPU Simulator

- General and extensible architecture
- Fast & accurate simulation

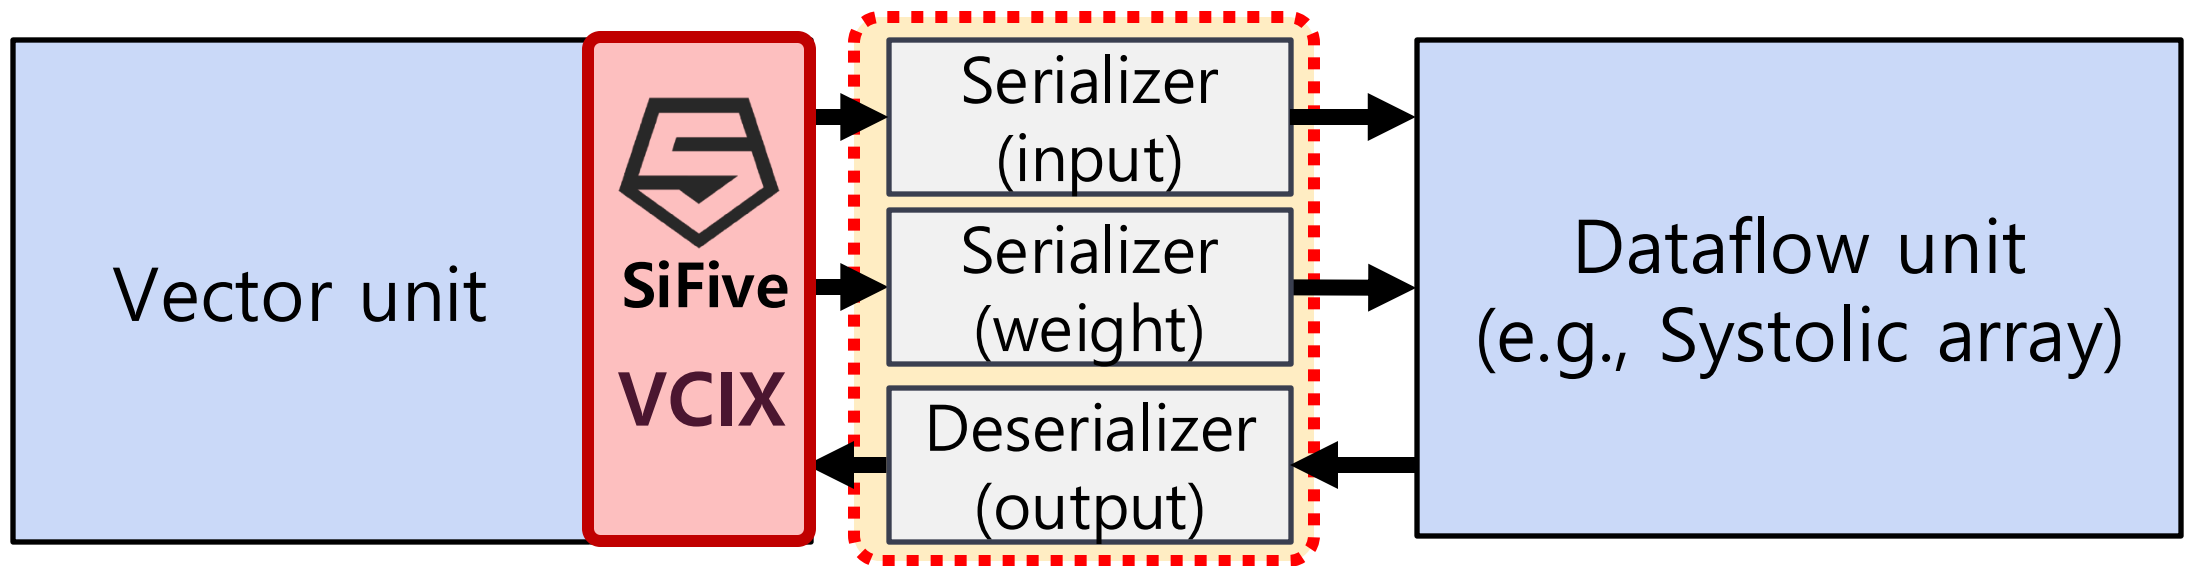
Modeling Common Building Blocks in NPU Cores



Extensible NPU ISA

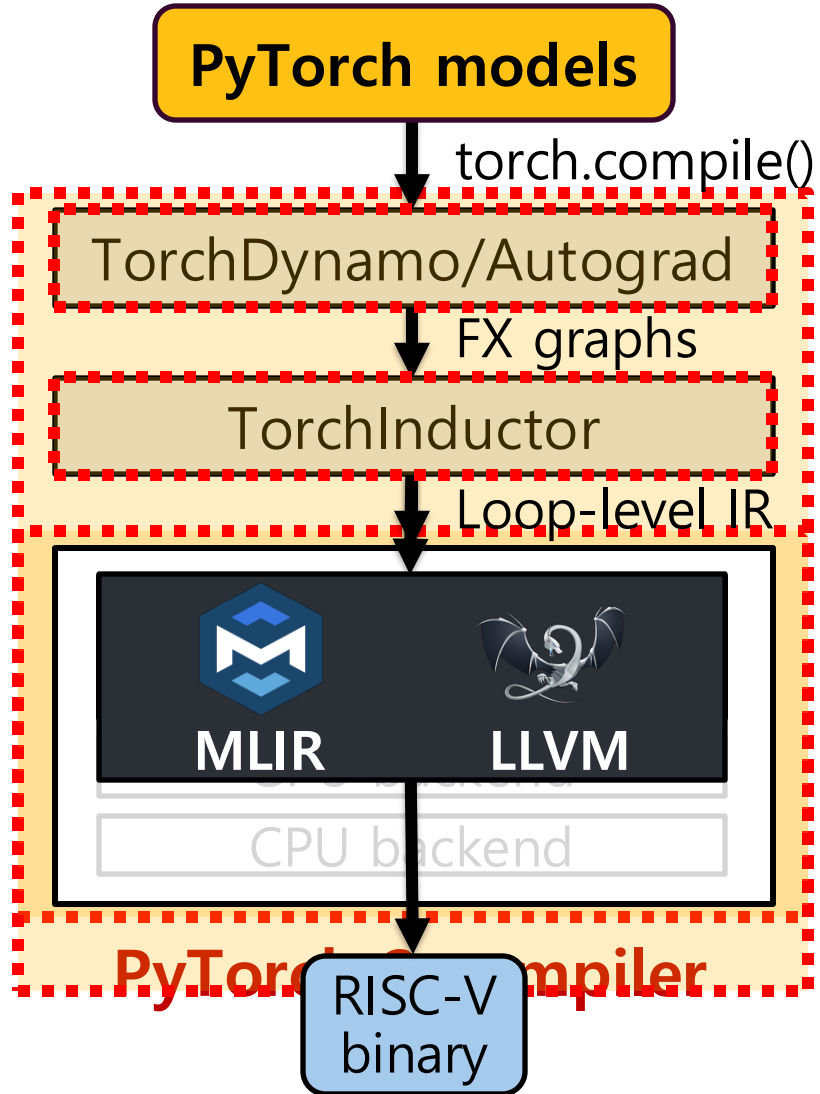
- **Scalar unit**
 - **Vector unit**
- } **RISC-V + Vector extension**
- **DMA engine**
 - **SFU.**
 - **Dataflow unit**
- } **Custom Instructions**

VCIX: A Generic Interface for Dataflow Units



We adopt VCIX, a **generic interface** for diverse dataflow units

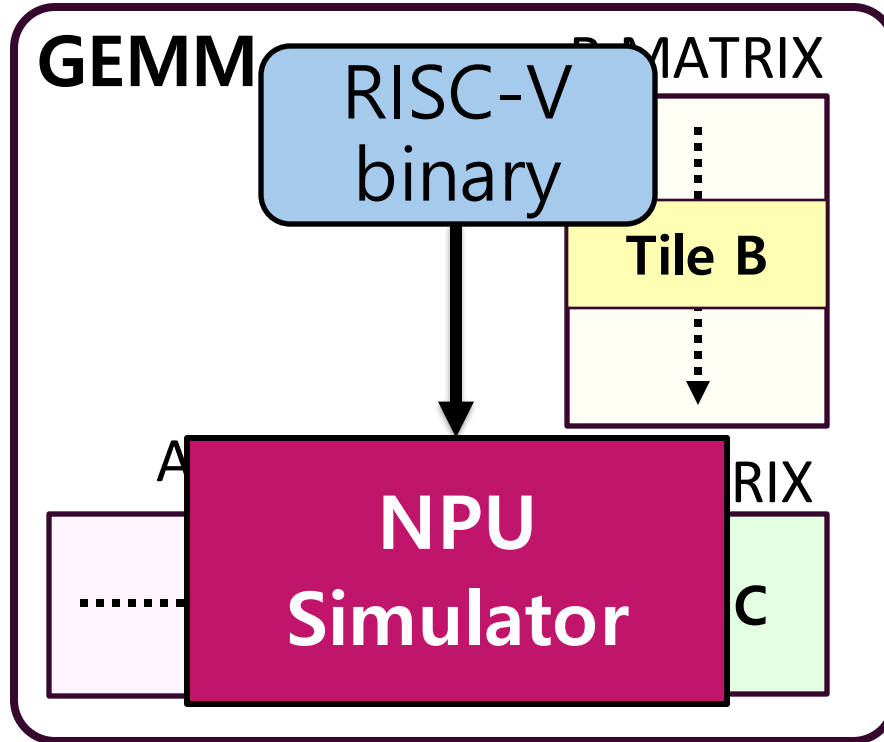
PyTorchSim Compilation Flow



- **Benefits of integration with PyTorch 2**
 - ▶ **Simulate existing PyTorch models** without any need for manual conversion
 - ▶ Supports **training simulation** through PyTorch's auto-differentiation
 - ▶ Leverage existing **target-independent optimizations**
- **NPU backend** generates a **RISC-V** binary
 - ▶ Applies target-dependent optimizations

Instruction-Level Simulation (ILS)

- Simulate every single instruction one-by-one
- Inherently slow**



RISC-V binary (GEMM)

**DMA load
Tile A**

```
.insn r 43, 3, 0, zero, t2, t4
.insn r 43, 3, 4, zero, a2, ra
.insn r 43, 3, 5, zero, a2, s8
.insn r 43, 3, 2, zero, a3, s0
```

**DMA load
Tile B**

```
add a0, s1, a6
.insn r 43, 3, 0, zero, a7, t4
.insn r 43, 3, 5, zero, a2, s8
.insn r 43, 3, 2, zero, a0, a5
```

```
addi a1, a1, 1
addi s1, s1, 512
addi a5, a5, 1024
bge s3, a4, .LBB0_16
j .LBB0_13
...
```


```
add a0, a0, s7
vmv1r.v v11, v6
```

**Matrix multiply
(~100 instructions)**

```
vle32.v v11, (a0), v0.t
sf.vc.iv 1, 0, v8, 0
sf.vc.iv 1, 0, v9, 0
...
sf.vc.iv 1, 0, v8, 0
sf.vc.iv 1, 0, v9, 0
```

```
li s11, 136
bltz s3, .LBB0_21
```

Key Insight behind NPU Execution

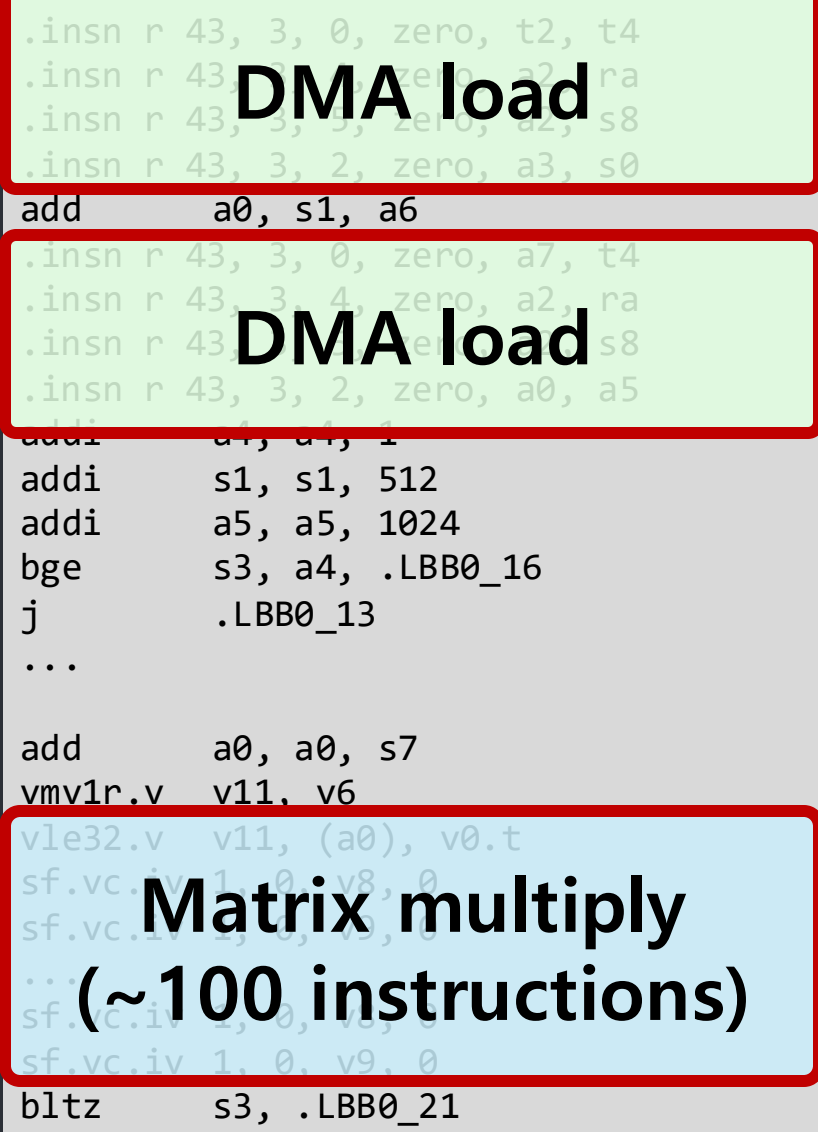


- Intra-core **compute latency** is *deterministic*


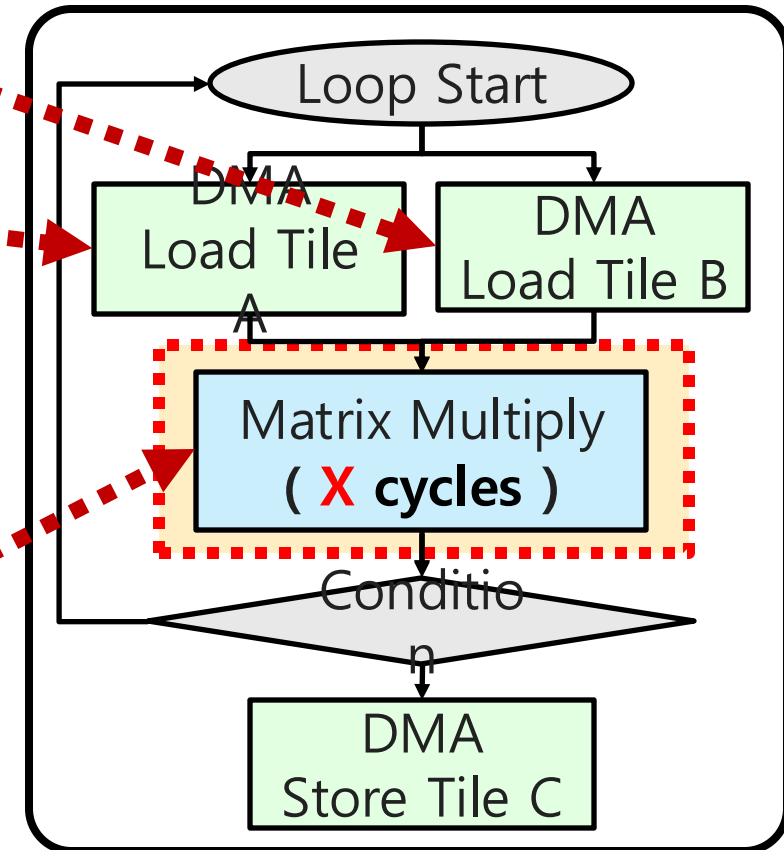
- **NoC & DRAM latencies** remain *non-deterministic* due to inter-core contention and traffic

Tile Operation Graph (TOG)

RISC-V binary (GEMM)



Tile Operation Graph (TOG)



Tile-Level Simulation (TLS) Execution Flow

Tile-Level Simulation works in a two-step flow

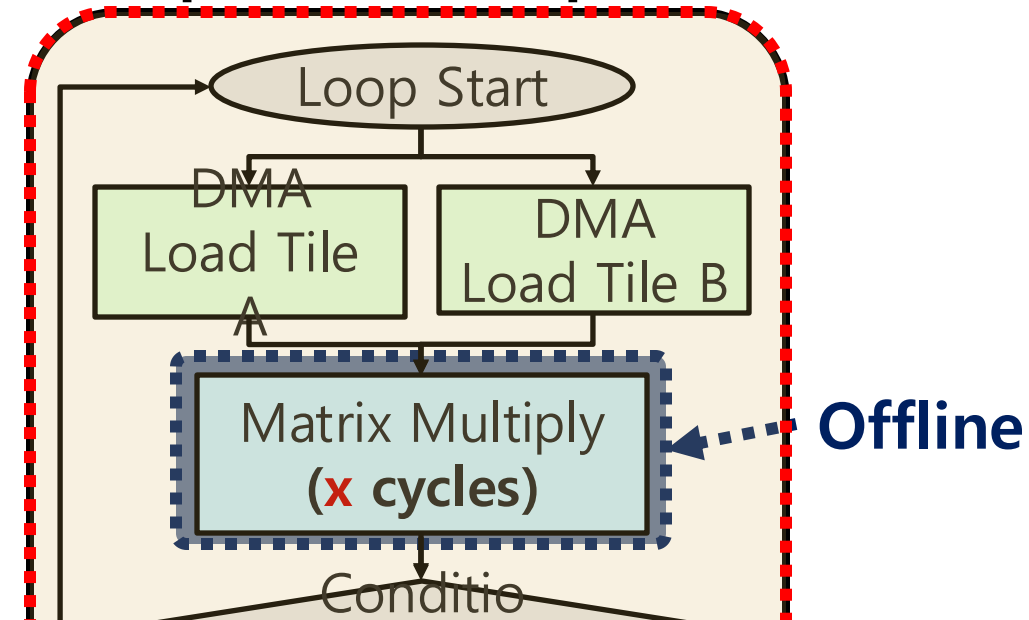
Offline (compile time)

- Obtain the deterministic **compute latency** in the TOG

Online (simulation time)

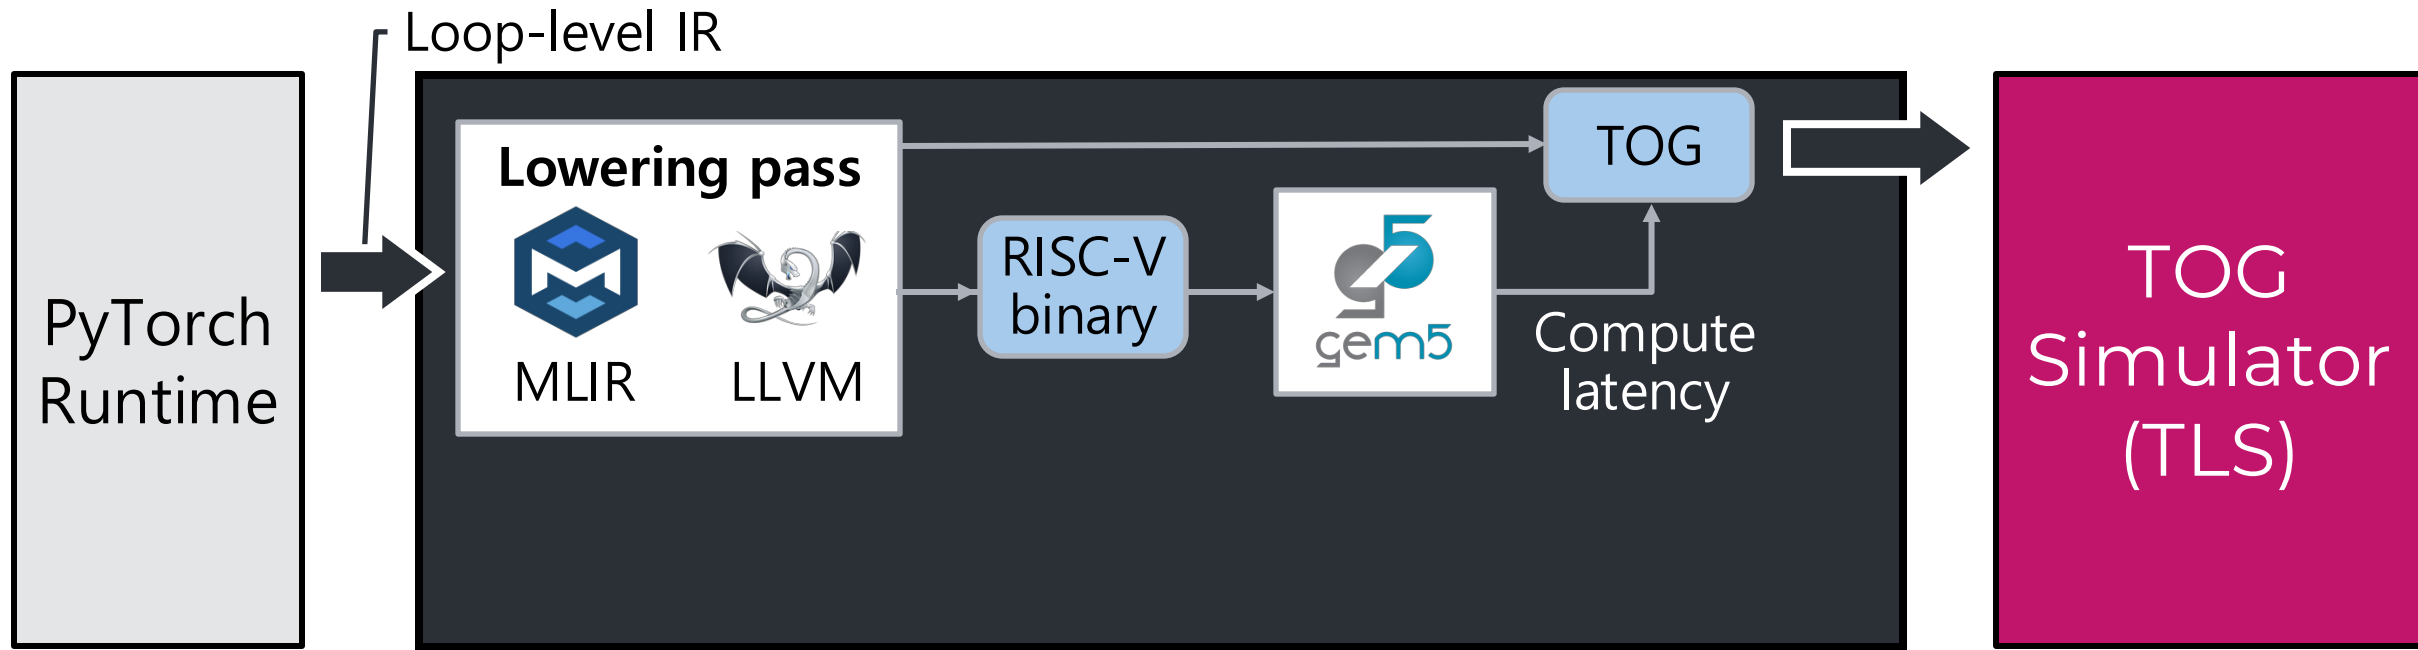
- Reuse the obtained compute latency

Tile Operation Graph (TOG)



TLS can achieve both **high simulation speed** and **accuracy**

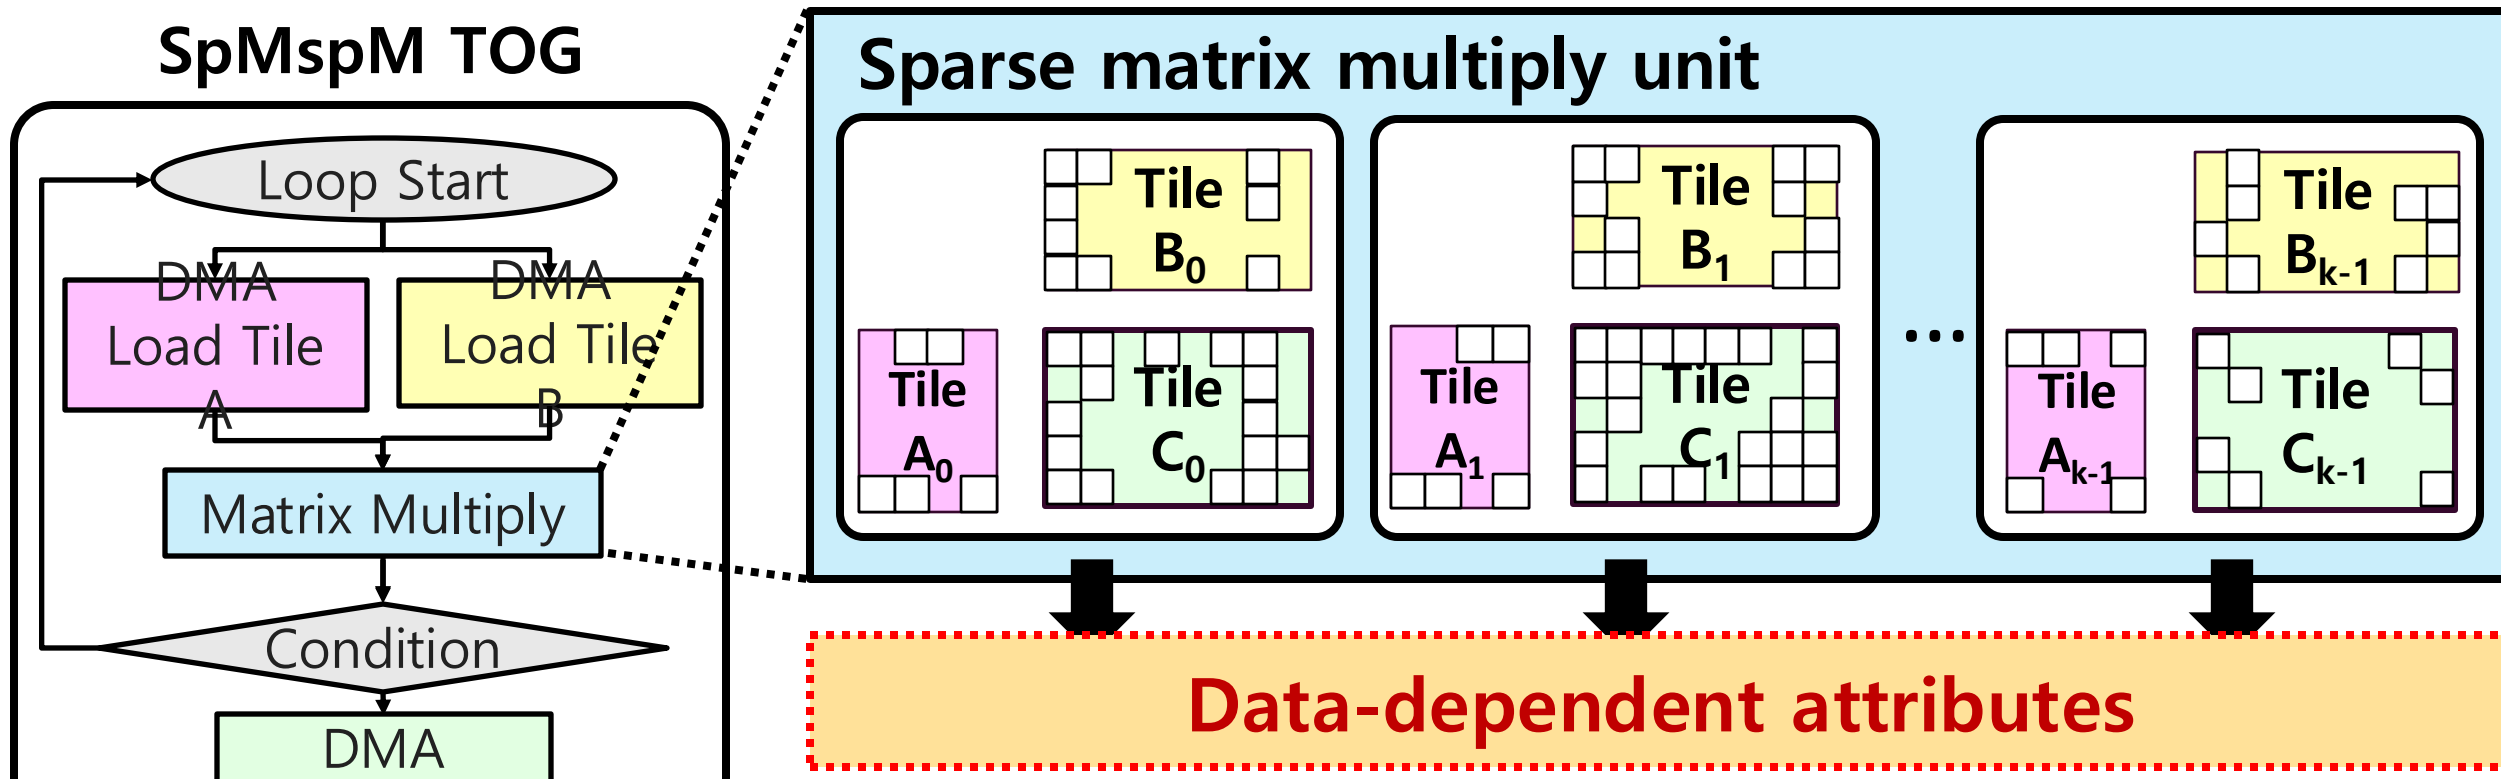
PyTorchSim TLS Compilation Flow



- **Gem5 (timing simulator)** models the compute latency for TLS

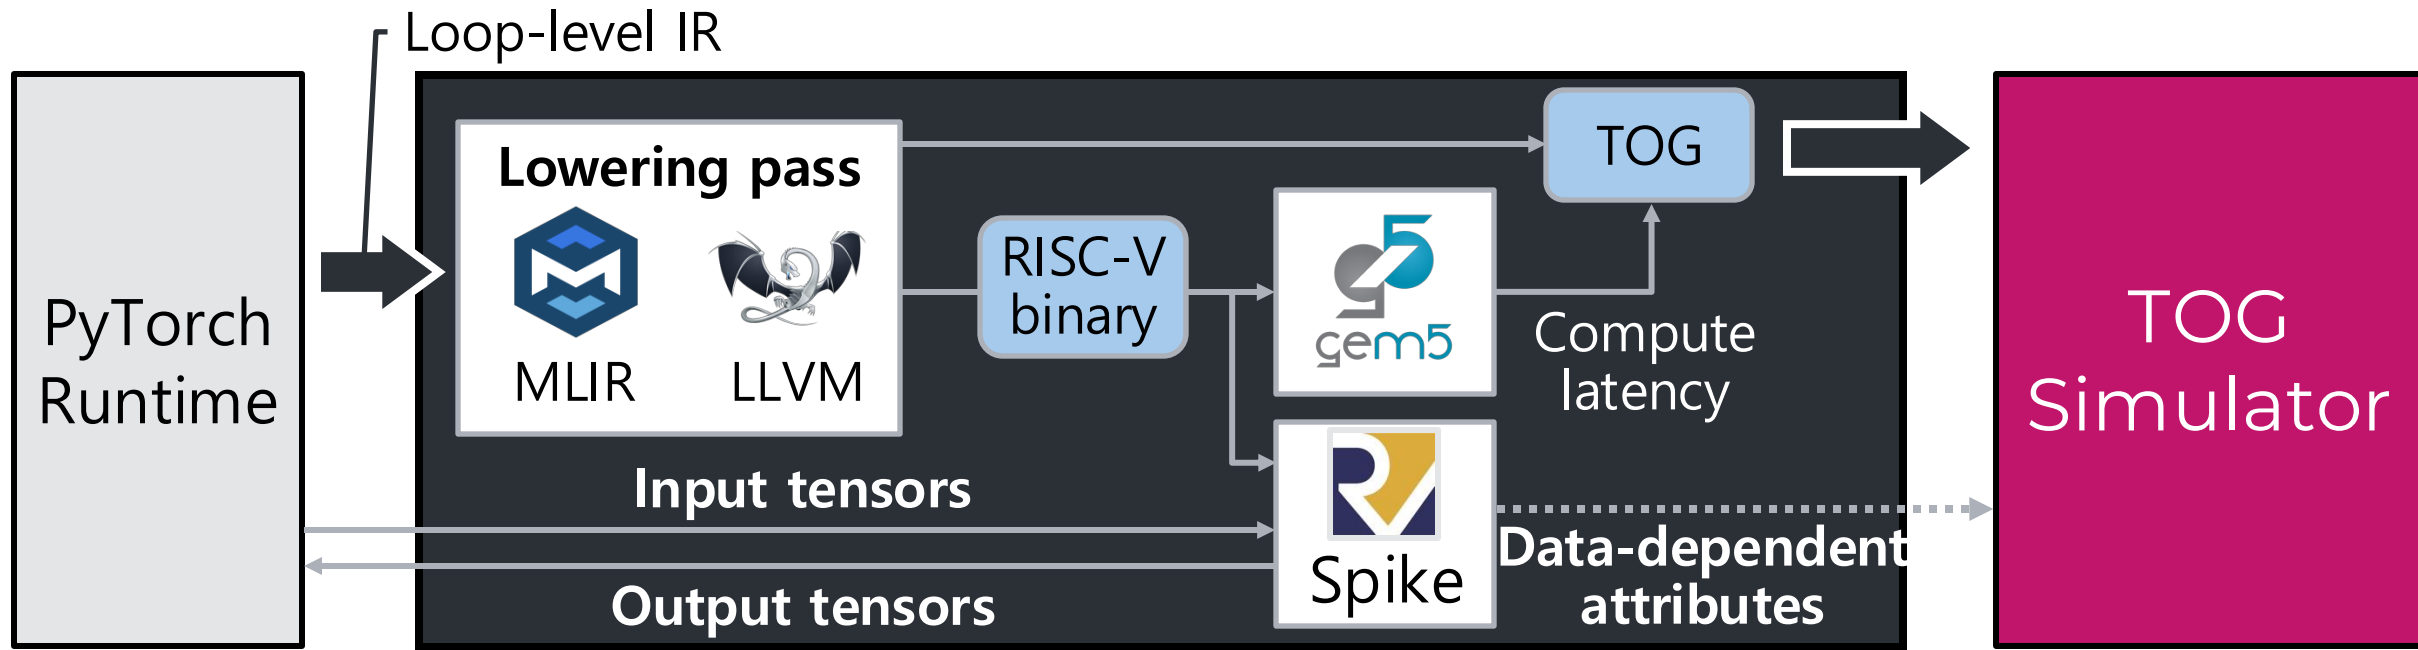
Challenge: Tile operation latency can depend on input data!

Data-Dependent Timing Behavior Example



For each **specific pair** of input tiles,
the compute latency is still **deterministic**

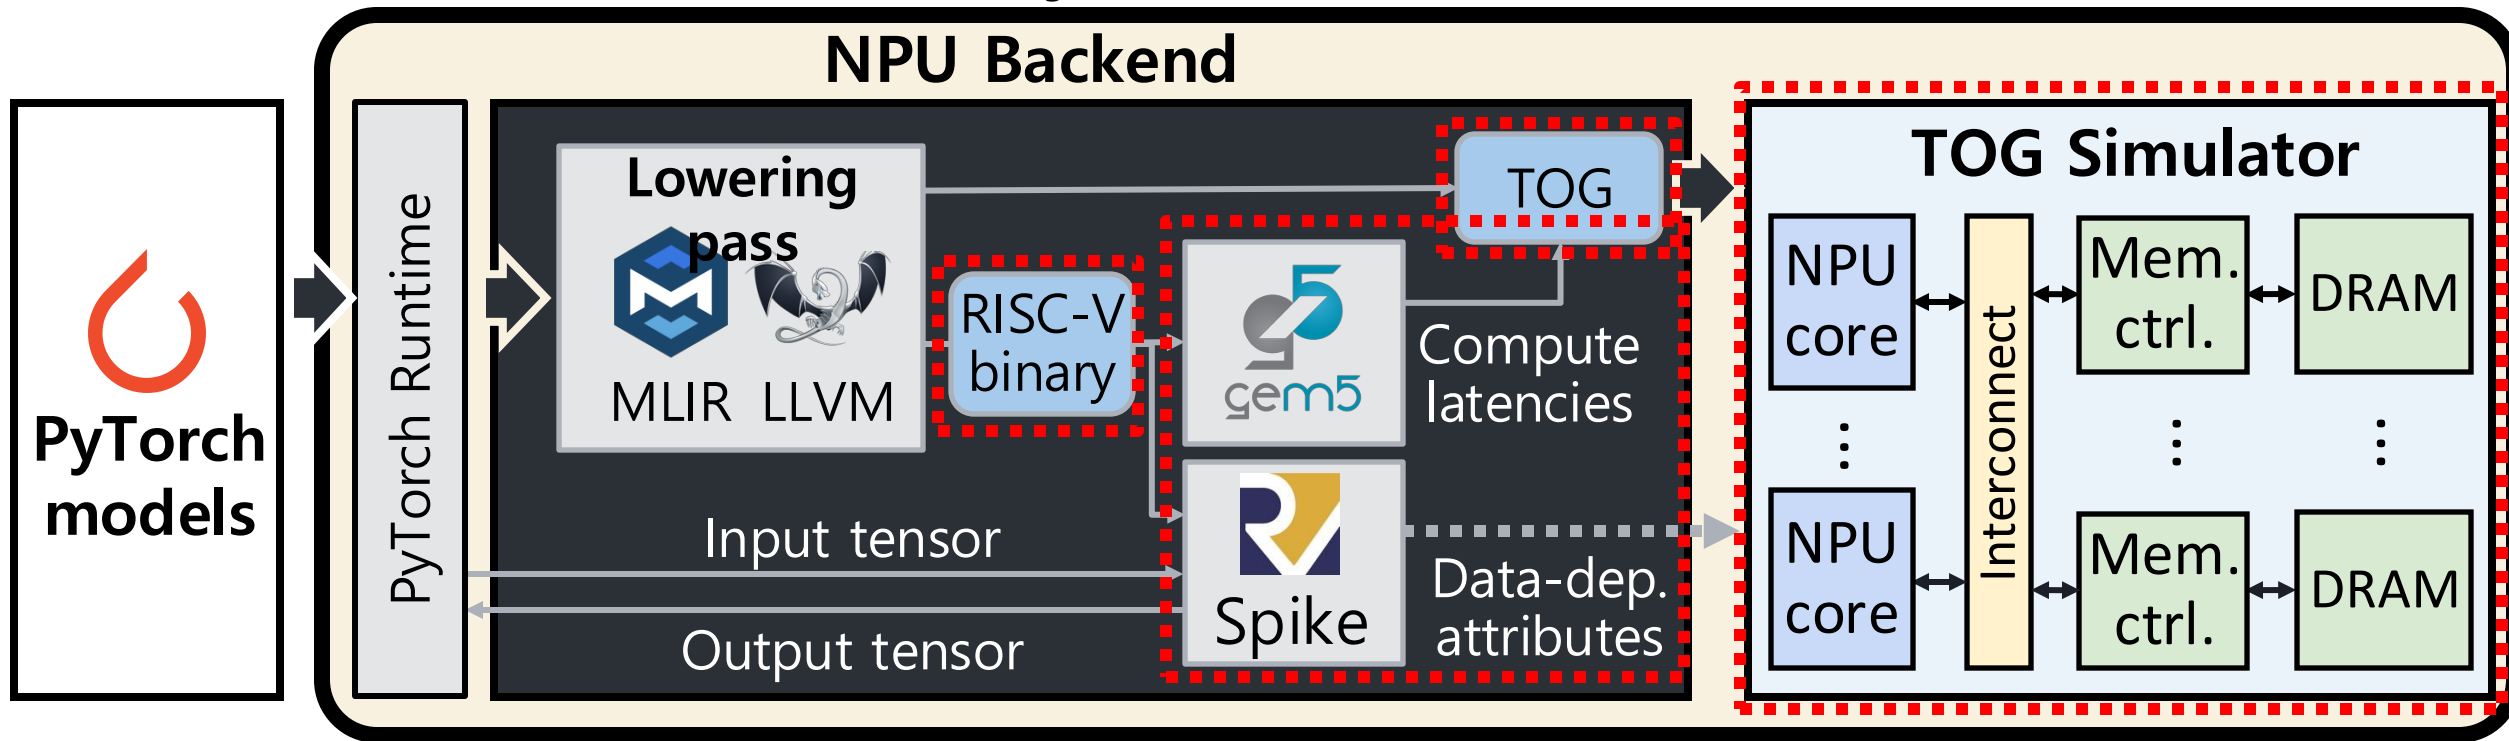
PyTorchSim TLS Compilation Flow



- **Spike (functional simulator)** executes the binary with PyTorch input:
 - To obtain data-dependent attributes
 - To validate correctness of compiled binary

Putting It All Together

PyTorchSim Framework



Evaluation Methodology

Accuracy validation

- **Target:** Real Google TPU v3
- **Baselines:** SCALE-Sim v3, mNPUSim, Timeloop, MAESTRO

	Cores & clock	Systolic arrays	Vector lanes (# of ALUs)	Scratchpad	DRAM
TPU v3	1 core @ 940 MHz	(128x128) x 2	128 (16 ALUs each)	32 MB	4 HBM2 (960 GB/s)


Simulation speed

- **Baselines:** Accel-Sim, mNPUSim
 - We selected Accel-Sim for its rich features and GPUs' dominance in deep learning

Workloads


- **Kernels:** GEMM, convolution, layer normalization, softmax, attention
- **Full models:** ResNet18/50, Bert-base/large

Evaluation: Validation against Real TPU v3



PyTorchSim can **accurately** simulate full models end-to-end

Evaluation: Speedup

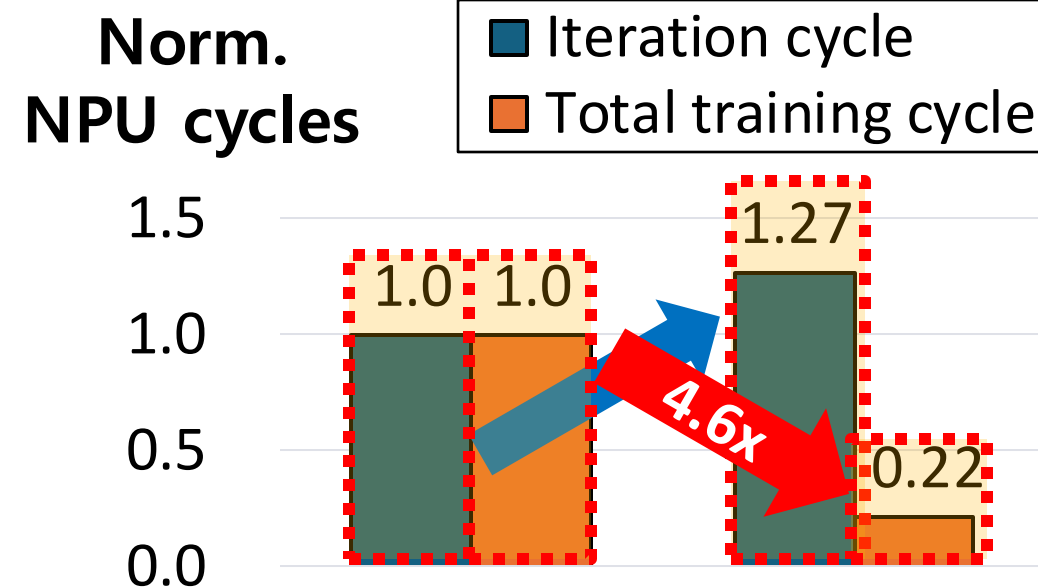



PyTorchSim achieves **significant simulation speedup**

Contents

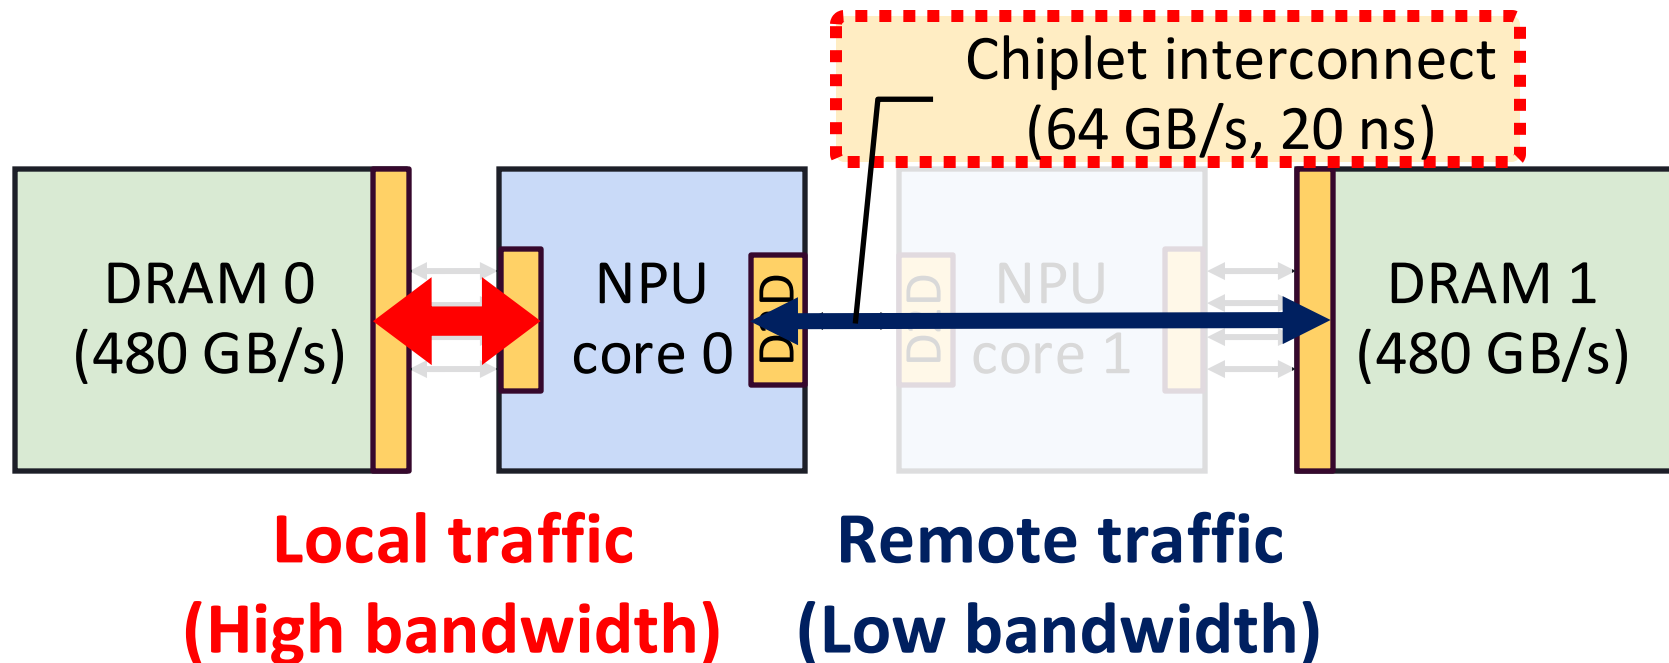
- Background / Motivation
- PyTorchSim
 - ▶ Overview
 - ▶ NPU core modeling
 - ▶ Compiler workflow
- Methodology & Evaluation
- **Case Studies (CS)**
 - ▶ Impact of DNN training hyperparameter
 - ▶ Impact of data placement for chiplet-based NPUs
- Summary

CS1: Impact of DNN Training Hyperparameter

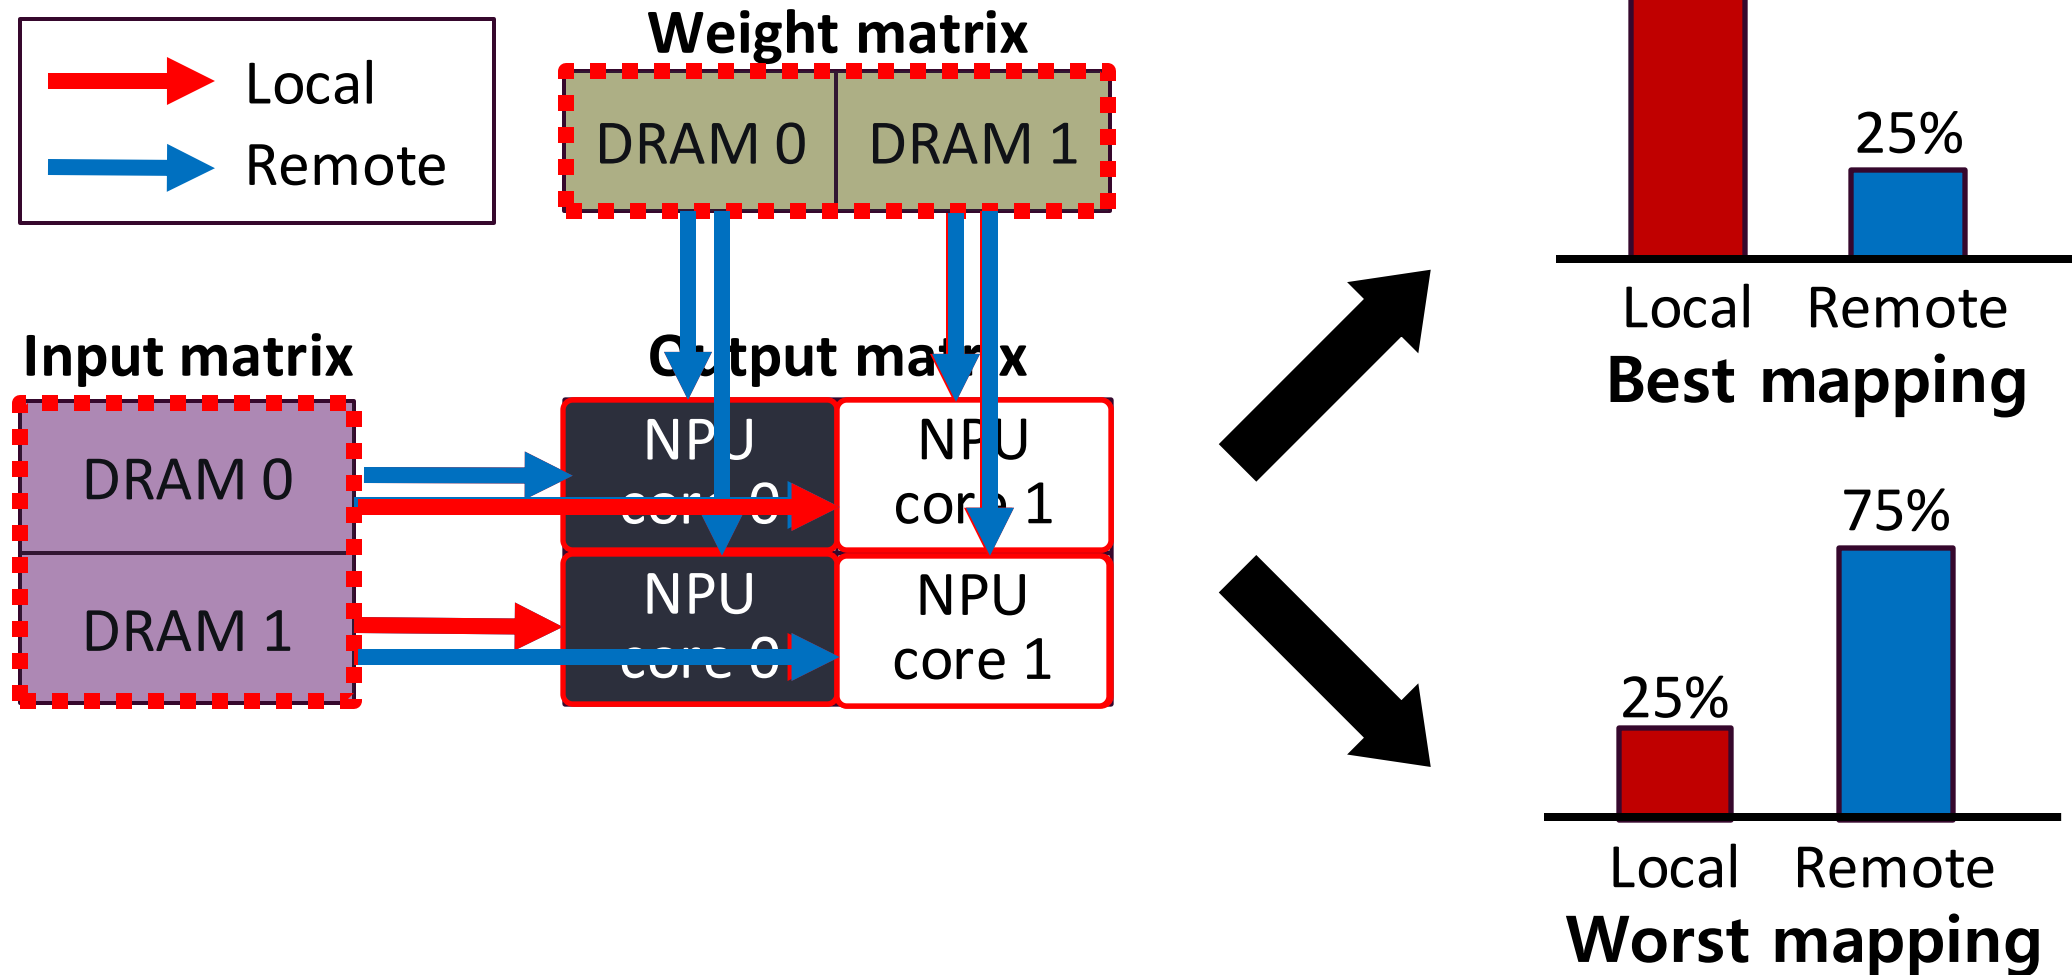


PyTorchSim enables studying **training** time behaviors of systems


CS2: Impact of Data Placement for Chiplet-based NPUs



CS2: Impact of Data Placement for Chiplet-based NPUs



CS2: Impact of Data Placement for Chiplet-based NPUs



- High remote traffic becomes the severe bottleneck

PyTorchSim enables the study of **diverse NPU architectures**

More Results and Discussions in the Paper

Case studies

- Heterogeneous dense-sparse NPU
- DNN inference with multi-model tenancy
- Compiler optimization impact

Discussions

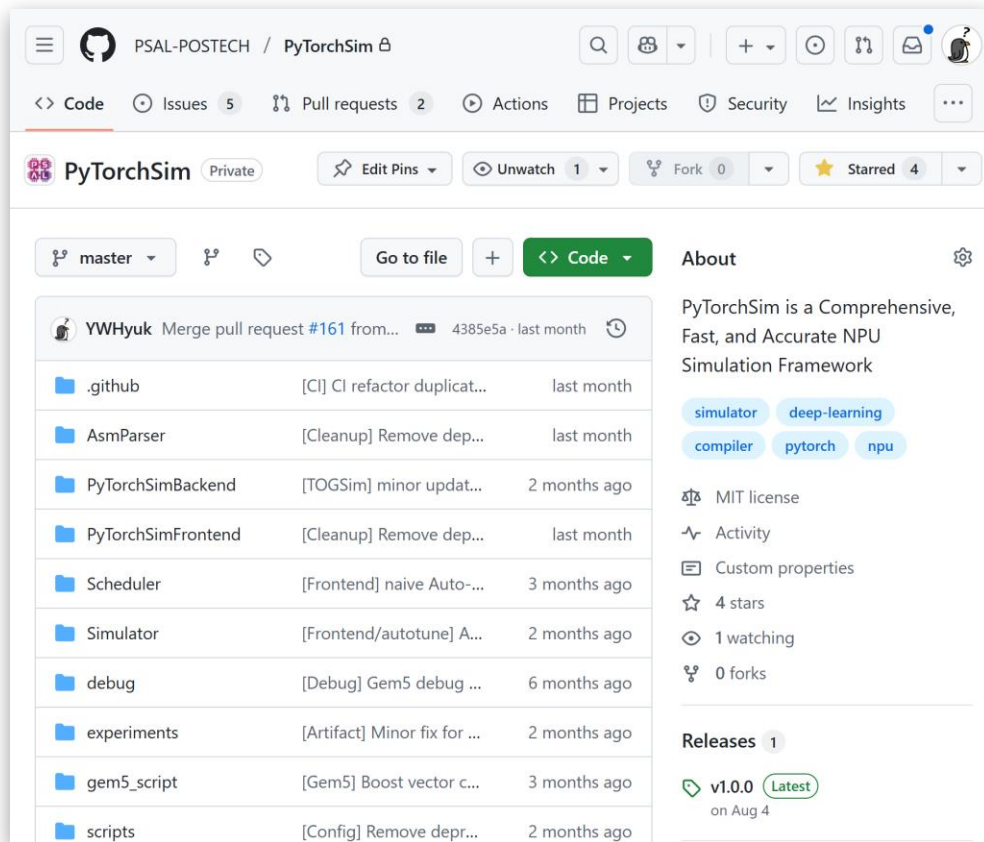
- Extension for other deep learning frameworks
- Extension for modeling GPUs and large-scale systems

Summary

- AI workloads require a comprehensive, fast, and accurate simulator
- To address this, we propose **PyTorchSim**
 - **High-speed, high-accuracy** through our **Tile-Level Simulation (TLS)**
 - **Integration with PyTorch 2 compilation flow**
 - **General and extensible architecture** through a **RISC-V-based NPU ISA**
- Compared to prior NPU simulators, PyTorchSim:
 - Enables various case studies not supported by prior simulators
 - Significantly improves **accuracy** and **simulation speed**

PyTorchSim is Open Source

- Contributions are welcome!



Model Zoo

Model	Source	Status	Note
ResNet-18			channel last format
ResNet-50			channel last format
BERT			
GPT-2			
ViT			
Mistral			
Diffusion			
Llama-4			Developing
DeepSeek v1			Developing

GitHub Link

